STATS/CSE 780
Assignment 3

Pao Zhu Vivian Hsu (Student Number: 400547994)

2023-11-07

	data_type	min	max	nulls_blanks
Transportation.expense	integer	118.000	388.000	0
Distance.from.Residence.to.Work	integer	5.000	52.000	0
Service.time	integer	1.000	29.000	0
Age	integer	27.000	58.000	0
Work.load.Average.day	numeric	205.917	378.884	0
Hit.target	integer	81.000	100.000	0
Son	integer	0.000	4.000	0
Pet	integer	0.000	8.000	0
Weight	integer	56.000	108.000	0
Height	integer	163.000	196.000	0
Body.mass.index	integer	19.000	38.000	0
Absenteeism.time.in.hours	integer	0.000	120.000	0

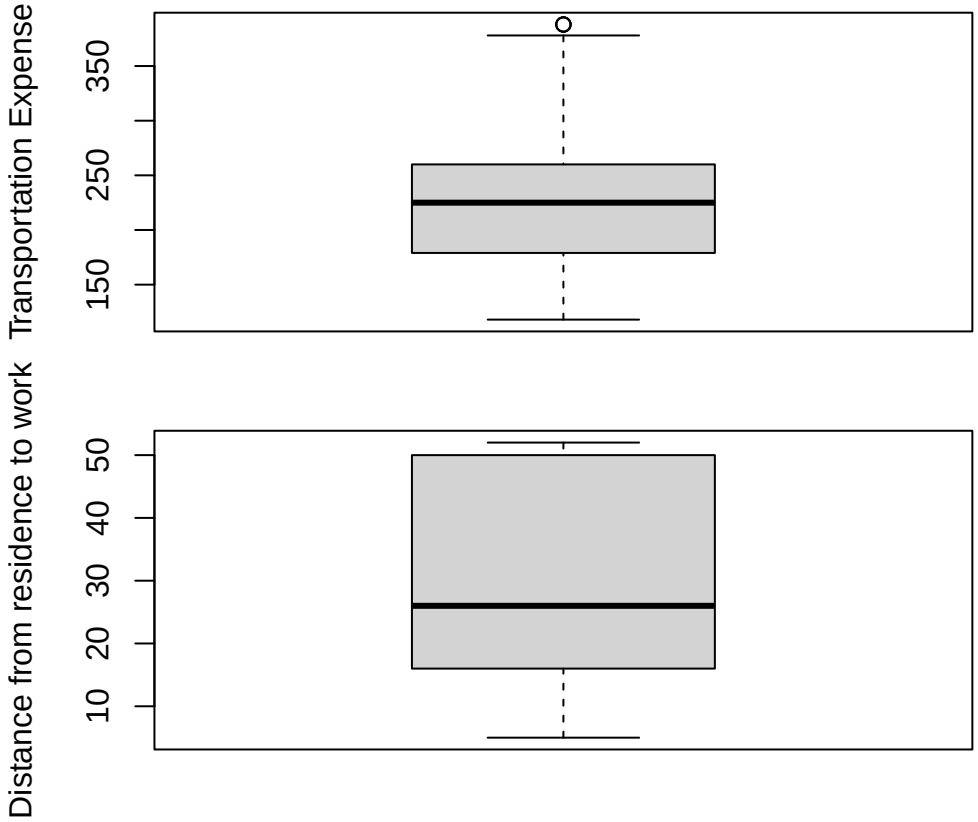

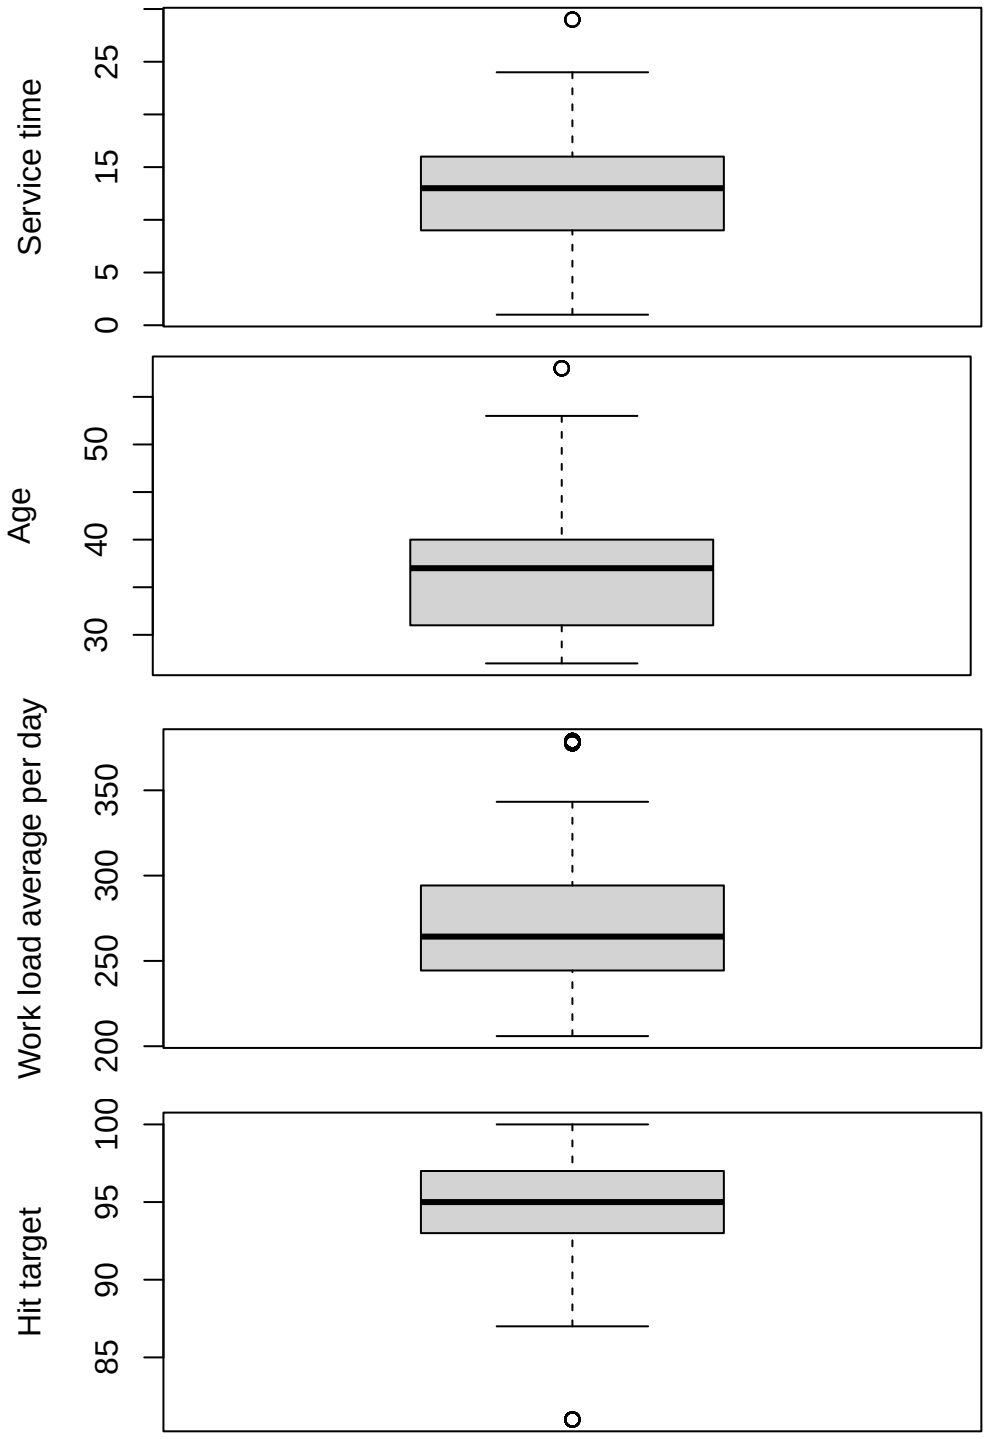
Service time
0
5
15
25
Age
30
40
50
Work load average per day
200
250
300
350
Hit target
85
90
95
100

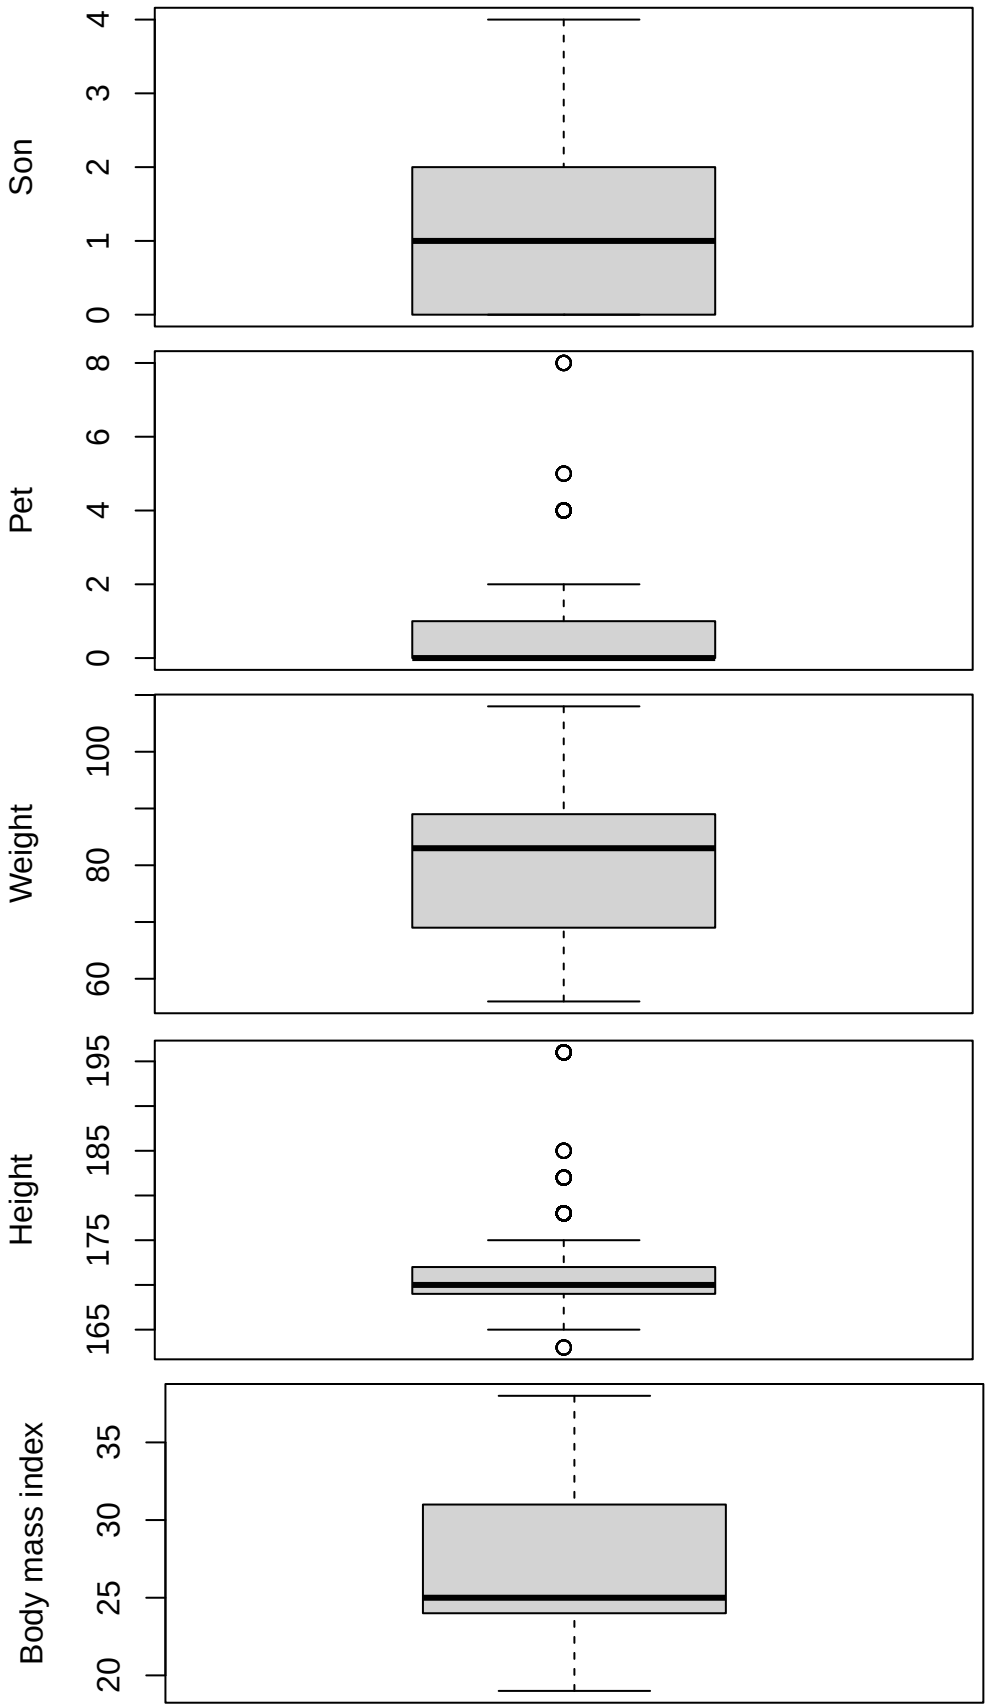
Son
0
1
2
3
4
Pet
0
2
4
6
8
Weight
60
80
100
Height
165
175
185
195
Body mass index
20
25
30
35

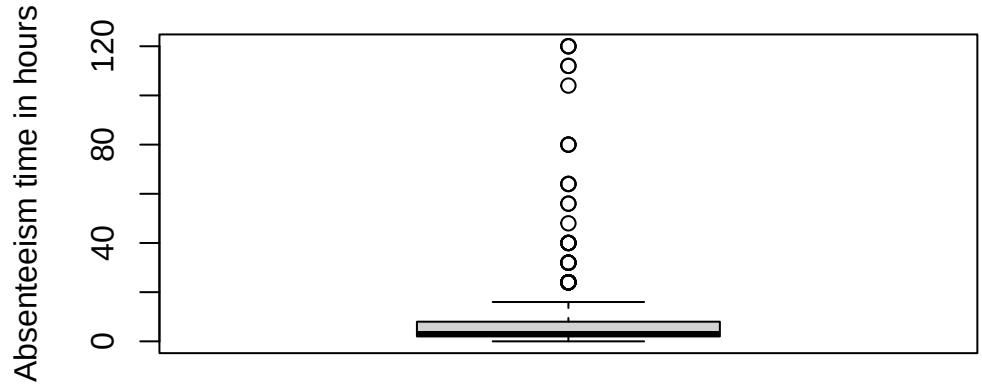

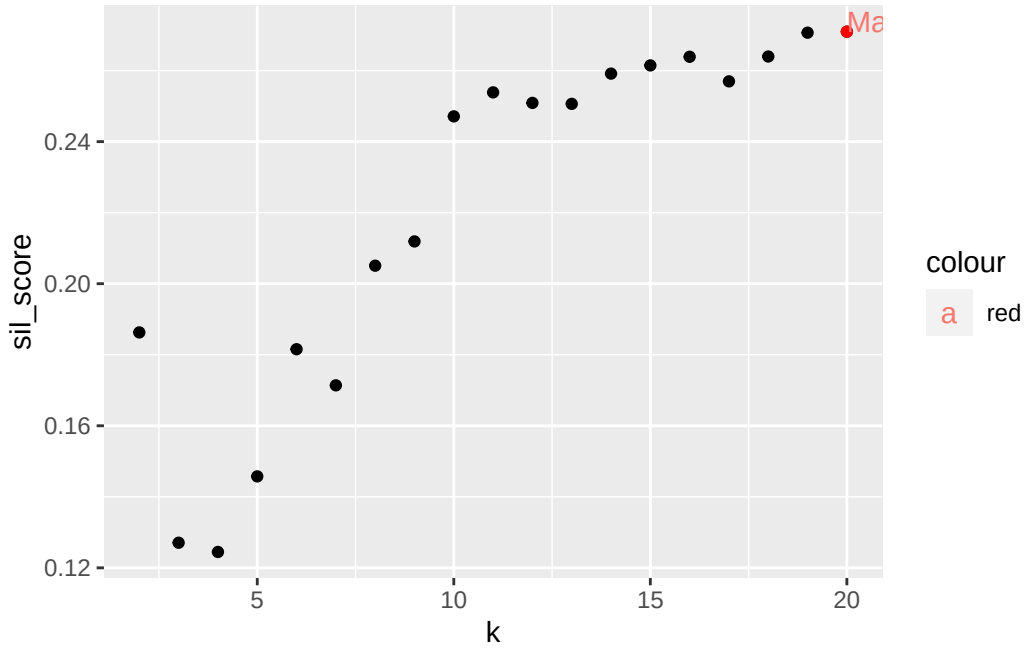

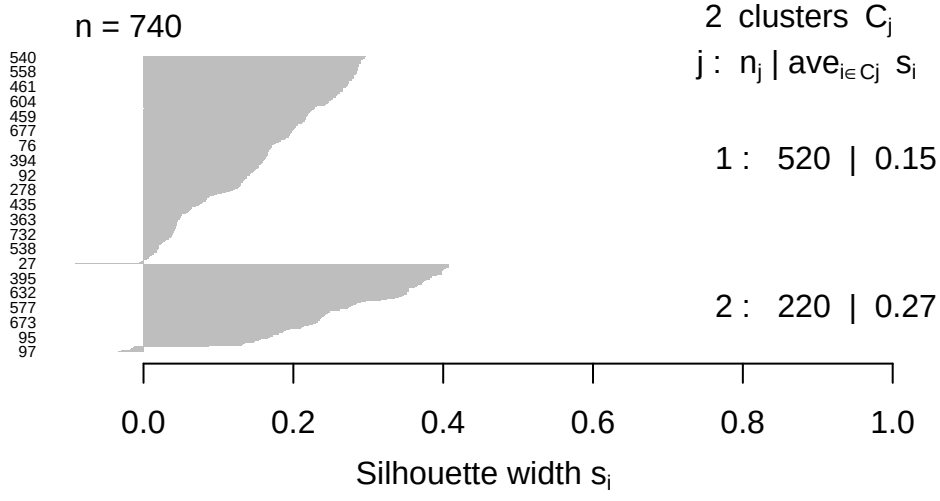

Average silhouette width : 0.19

```
[1] 0.8410708
```

```
[1] 0.09340035
```

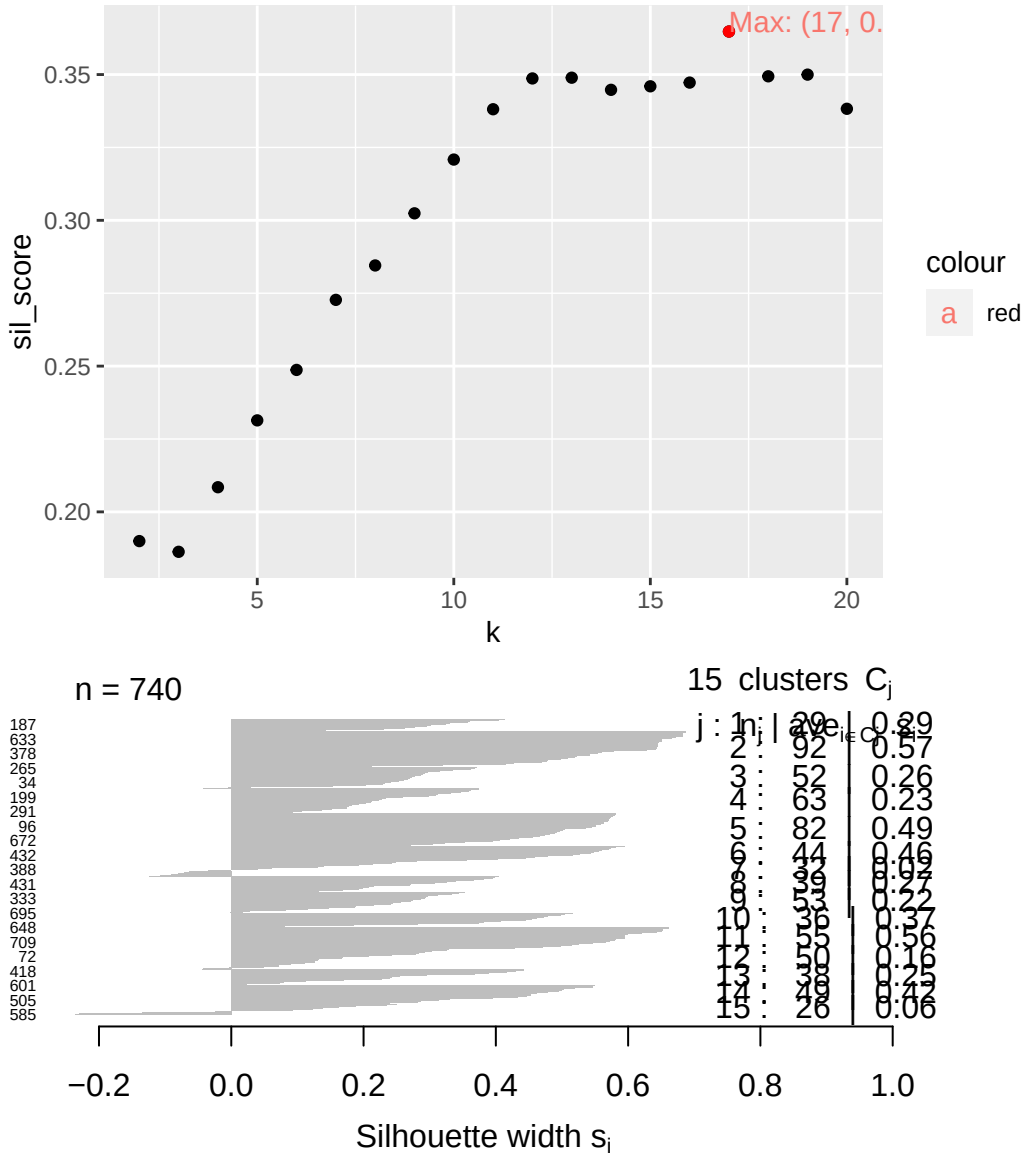
Max: (17, 0.
sil_score
k
colour
a
red
n = 740
15 clusters C_j
Silhouette width s_i
Average silhouette width : 0.35

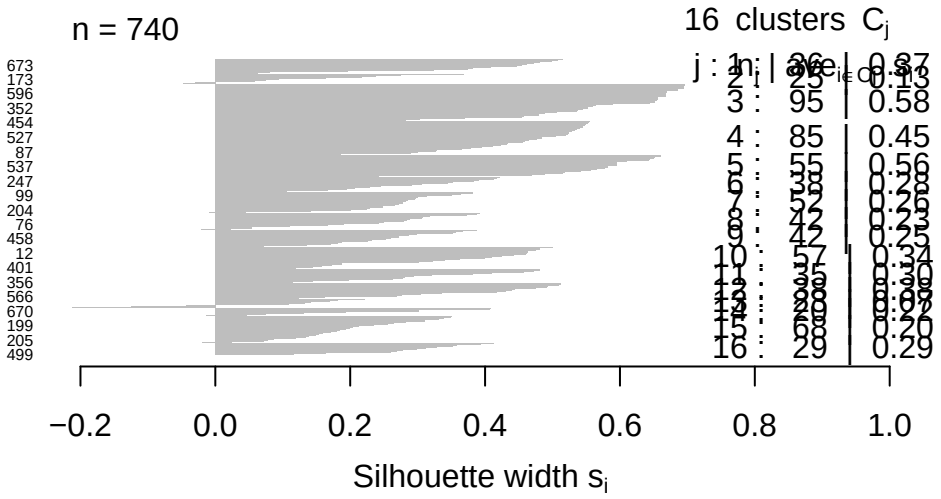

Average silhouette width : 0.35

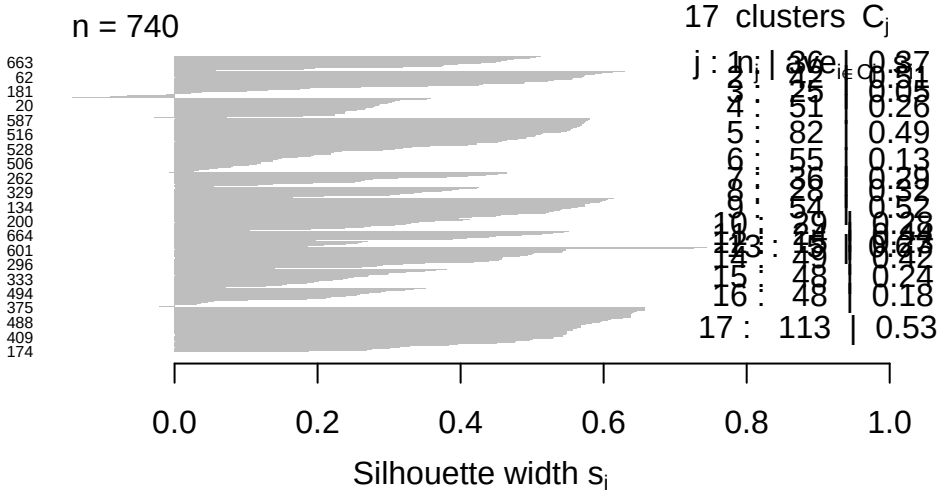

Average silhouette width : 0.36

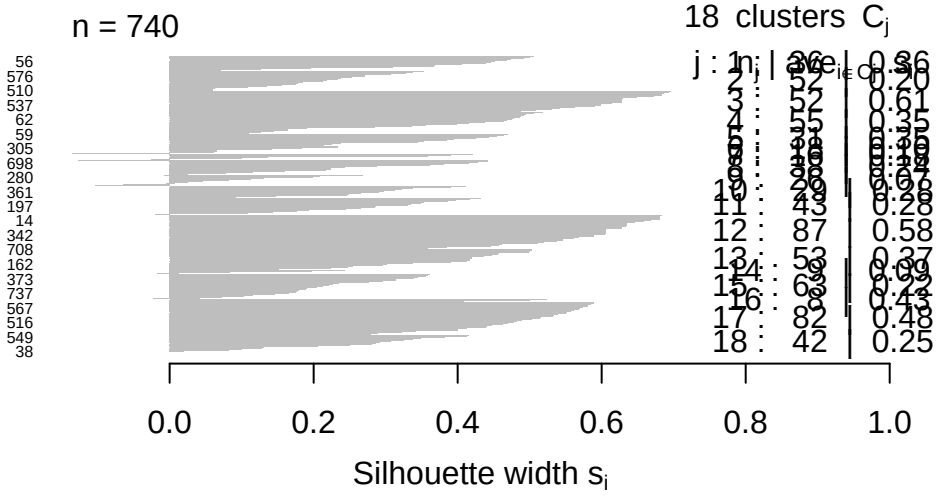

Average silhouette width : 0.35

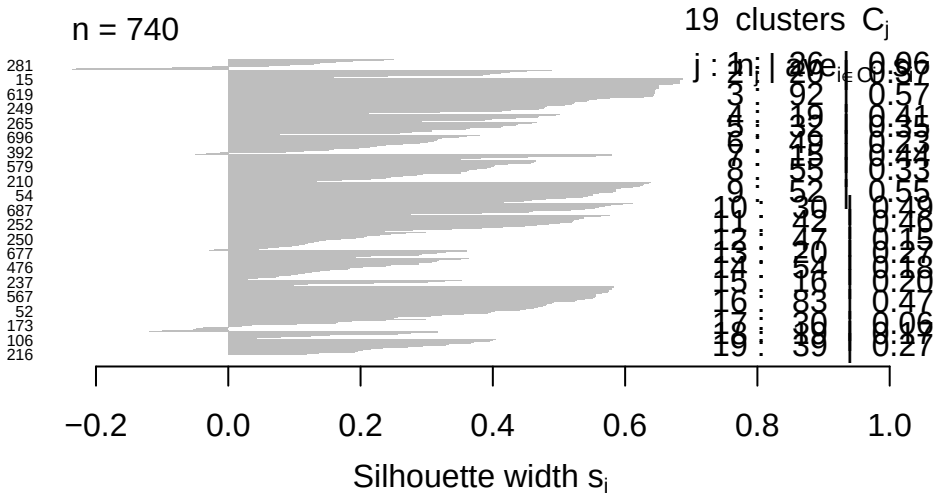

Average silhouette width : 0.35

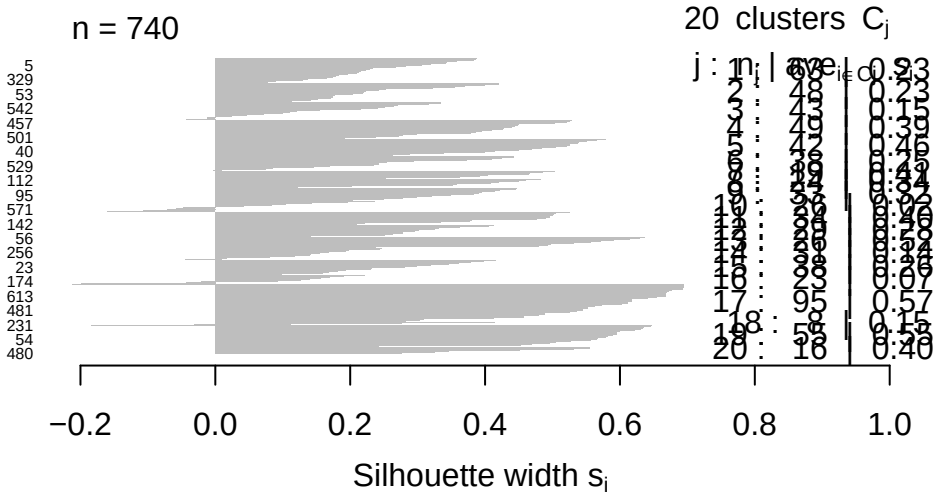

Average silhouette width : 0.34

```
[1] 0.8539041
```

```
[1] 0.1017307
```

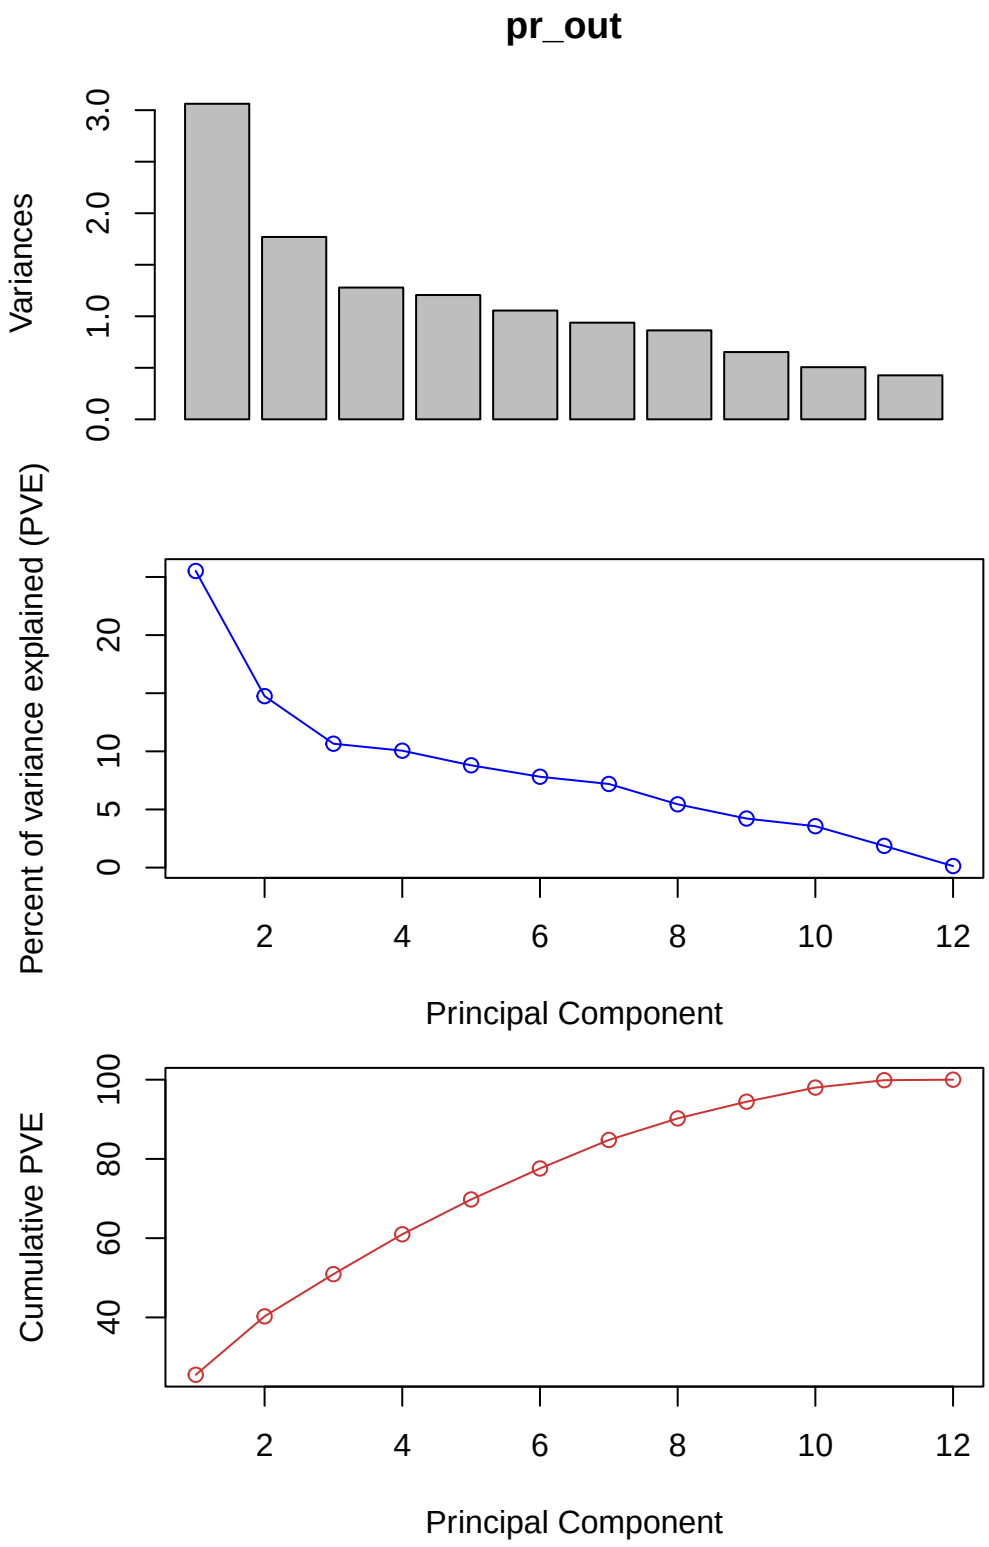
pr_out
Variances
0.0
1.0
2.0
3.0
Percent of variance explained (PVE)
0
5
10
20
2
4
6
8
10
12
Principal Component
Cumulative PVE
40
60
80
100
2
4
6
8
10
12
Principal Component

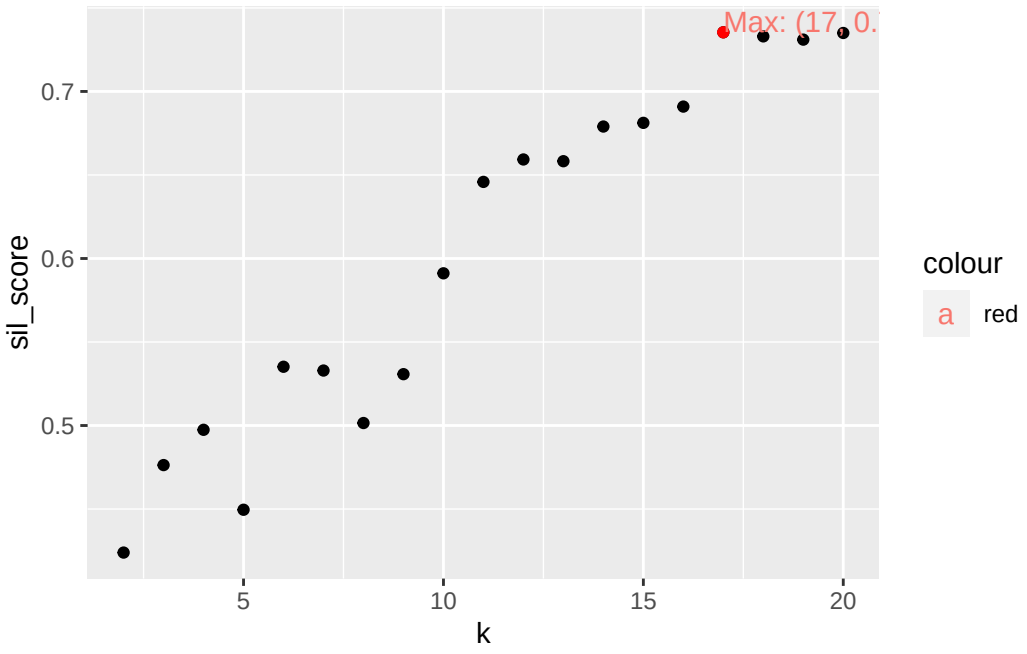

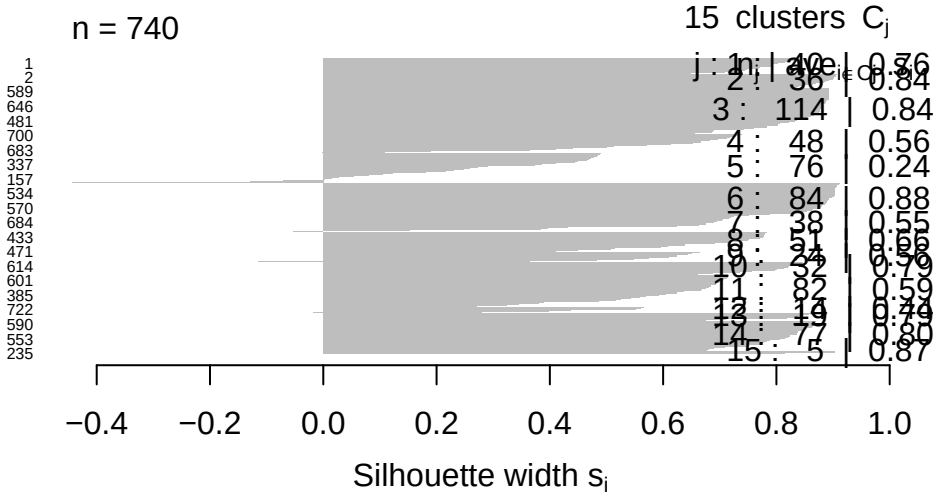

Average silhouette width : 0.68

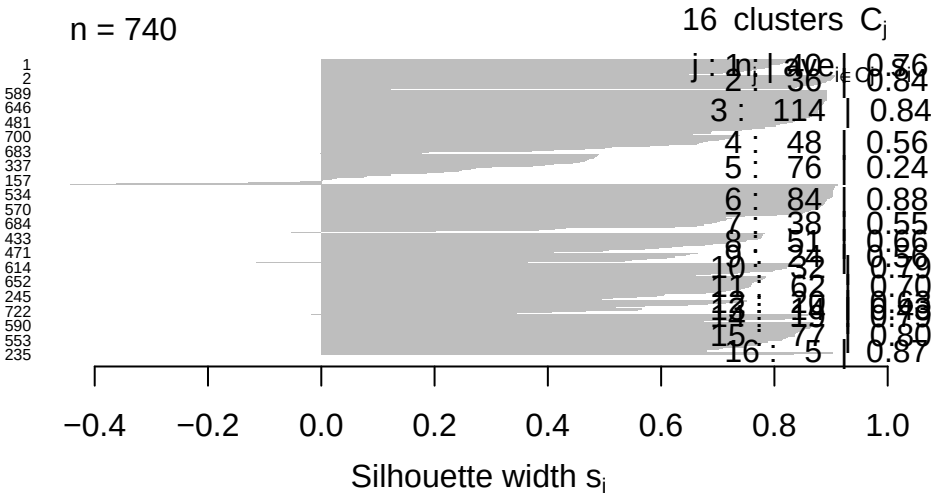

Average silhouette width : 0.69

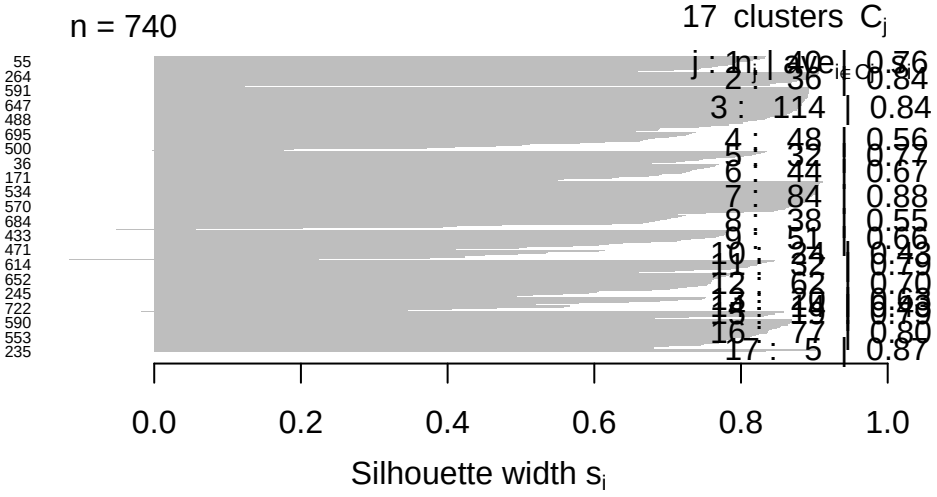

Average silhouette width : 0.74

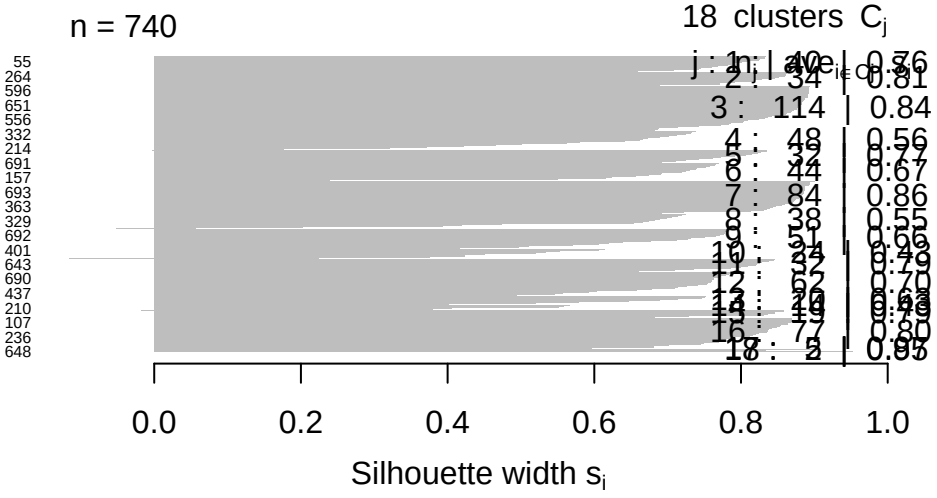

Average silhouette width : 0.73

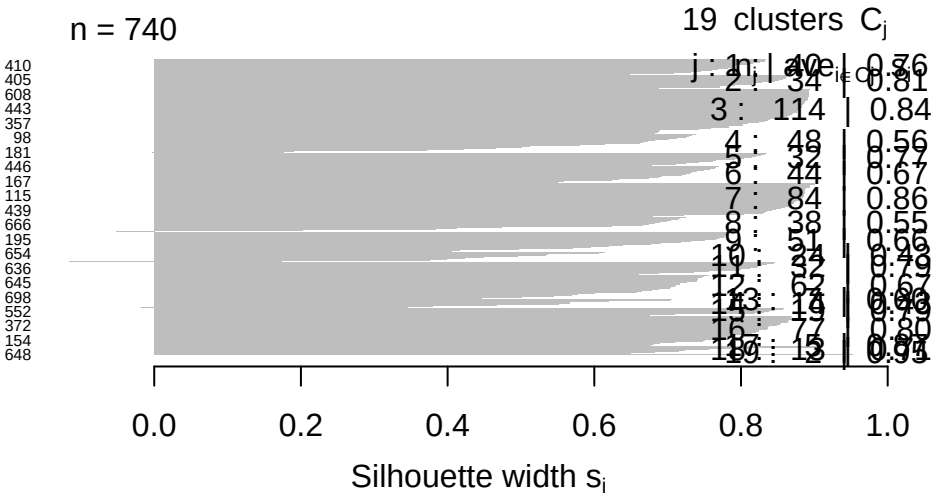

Average silhouette width : 0.73

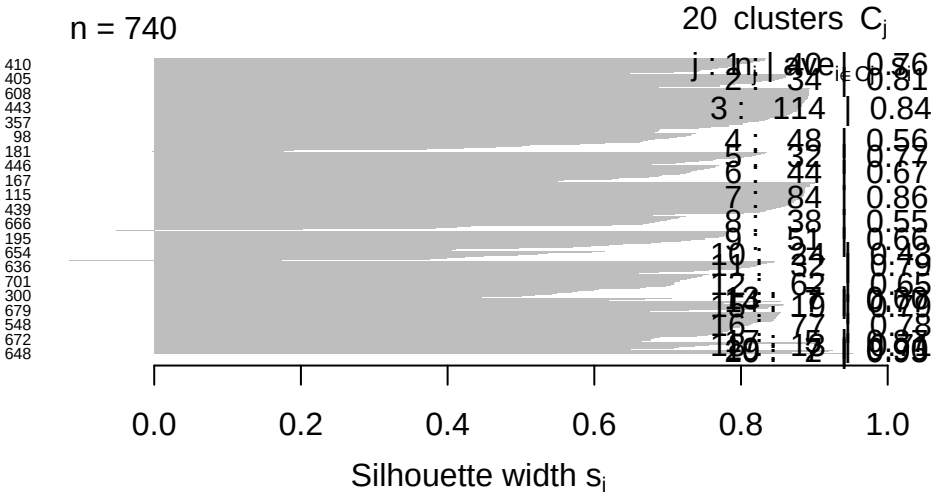

Average silhouette width : 0.74

```
[1] 0.8454778
```

```
[1] 0.1130238
```

Introduction

Methods

Discussion

Supplementary material

```
# ----- SETUP ----- #
# Load packages
packages <- c("knitr", "tidyverse", "ggplot2", "cluster", "fossil")
lapply(packages, library, character.only = TRUE)

# Read data, extract labels, and keep only quantitative data
absentData_raw <- read.csv("Absenteeism_at_work.csv", sep = ";")
absentData_lab <- absentData_raw$`Reason.for.absence`
absentData <- absentData_raw %>%
  select(-c("Reason.for.absence","ID","Month.of.absence","Day.of.the.week","Seasons",
            "Disciplinary.failure","Education","Social.drinker","Social.smoker"))
# ----- DATA EXPLORATION ----- #
# Check data types, min, max, and missing data
data_type <- sapply(absentData,class)
min <- sapply(absentData, function(col){min(col,na.rm=TRUE)})
max <- sapply(absentData, function(col){max(col,na.rm=TRUE)})
nulls <- sapply(absentData, function(col){sum(is.na(col))})
blanks <- sapply(absentData,
                 function(col){ifelse(is.na(sum(col == "")), 0, sum(col == ""))})
data_summary <- data.frame(row.names = names(nulls), data_type=data_type,
                           min=min, max=max, nulls_blanks=nulls+blanks)
kable(data_summary)

# Create box plots to check for outliers
b01 <- boxplot(absentData$Transportation.expense, ylab = "Transportation Expense")
b02 <- boxplot(absentData$Distance.from.Residence.to.Work,
               ylab = "Distance from residence to work")
b03 <- boxplot(absentData$Service.time, ylab = "Service time")
b04 <- boxplot(absentData$Age, ylab = "Age")
b05 <- boxplot(absentData$Work.load.Average.day, ylab = "Work load average per day")
```

```r
b06 <- boxplot(absentData$Hit.target, ylab = "Hit target")
b07 <- boxplot(absentData$Son, ylab = "Son")
b08 <- boxplot(absentData$Pet, ylab = "Pet")
b09 <- boxplot(absentData$Weight, ylab = "Weight")
b10 <- boxplot(absentData$Height, ylab = "Height")
b11 <- boxplot(absentData$Body.mass.index, ylab = "Body mass index")
b12 <- boxplot(absentData$Absenteeism.time.in.hours, ylab = "Absenteeism time in hours")

# ----- DATA CLEANSING -----
# Handle outliers by capping them using interquartile range
cap <- function(val, bplot) {
  lower_fence <- bplot$stats[2]-(1.5*(bplot$stats[4]-bplot$stats[2])) #Q1-1.5*IQR
  upper_fence <- bplot$stats[4]+(1.5*(bplot$stats[4]-bplot$stats[2])) #Q3+1.5*IQR
  val <- ifelse(val < lower_fence, lower_fence, val)
  val <- ifelse(val > upper_fence, upper_fence, val)
  val
}
absentData <- absentData %>%
  mutate(Transportation.expense = cap(val=Transportation.expense, bplot=b01),
         Service.time = cap(val=Service.time, bplot=b03),
         Age = cap(val=Age, bplot=b04),
         Work.load.Average.day = cap(val=Work.load.Average.day, bplot=b05),
         Hit.target = cap(val=Hit.target, bplot=b06),
         Pet = cap(val=Pet, bplot=b08),
         Height = cap(val=Height, bplot=b10),
         Absenteeism.time.in.hours = cap(val=Absenteeism.time.in.hours, bplot=b12))

# Scale the data
absentData_sd <- scale(absentData)

# ----- CLUSTERING FUNCTIONS ----- #
```

```r
# Get silhouette for k-means clustering
kmcSilK <- function(k, data){
  x_k <- kmeans(data, k, nstart = 20)
  silhouette(x_k$cluster, dist(data))
}

# Get silhouette for hierarchical clustering
hcSilK <- function(k, data){
  hc_out <- hclust(dist(data))
  hc_clusters <- cutree(hc_out, k)
  silhouette(hc_clusters, dist(data))
}

# Plot silhouette
plotSil <- function(sil){
  plot(sil, nmax= 800, cex.names=0.5, main = "", border=NA)
}

# Choose k using goodness-of-clustering
# k = the k values to test
# silFun = the silhouette function
# data = the data used in silFun
chooseK <- function(k, silFun, data) {

  # Get silhouettes and their widths
  sil_k <- lapply(k, silFun, data=data)
  sil_score <- sapply(sil_k, function(x) {mean(x[,"sil_width"])})

  # Find the k with the max width
  sil_max <- max(sil_score)
  sil_max_k <- match(sil_max, sil_score)+min(k)-1
```

```
  # Plot the silhouette widths and label the maximum
  silData <- tibble(k, sil_score)
  max_point <-tibble(k=sil_max_k,sil_score=sil_max)
  max_lab <- paste0("Max: (",sil_max_k,", ",round(sil_max,2), ")")
  plot <- ggplot(silData, aes(x=k, y=sil_score)) + geom_point() +
    geom_point(data=max_point, colour="red") +
    geom_text(data=max_point, aes(label=ifelse(k==sil_max_k,max_lab,"")), color="red"),hjust=

  # Return plot, silhouettes, and max k
  list(plot=plot, sil_k=sil_k, max_k=sil_max_k)
}

#plotSil(hcSilK(2, absentData_sd)) # hierarchical clustering
#plotSil(hcSilK(2, pr_out$x[, 1:2])) # hierarchical clustering after PCA
#chooseK(2:30, hcSilK, absentData_sd)
#chooseK(2:30, hcSilK, pr_out$x[, 1:2])

# ----- AGGLOMERATIVE HIERARCHICAL CLUSTERING ----- #
# Get silhouette scores for multiple k values and plot them
set.seed(3)
k <- c(2:20)
good_of_cluster <- chooseK(k, hcSilK, absentData_sd) # uses complete linkage
good_of_cluster$plot
plotSil(good_of_cluster$sil_k[[1]])
k <- good_of_cluster$max_k

# Perform hierarchical clustering using best k
set.seed(3)
hc_out <- hclust(dist(absentData_sd))
rand.index(cutree(hc_out, k), as.numeric(as.factor(absentData_lab)))
```

```r
adj.rand.index(cutree(hc_out, k), as.numeric(as.factor(absentData_lab)))

# ----- K-MEANS CLUSTERING ----- #
# Get silhouette scores for multiple k values and plot them
set.seed(3)
k <- c(2:20)
good_of_cluster <- chooseK(k, kmcSilK, absentData_sd)
good_of_cluster$plot

# Based off silhouette plots, choose the best k
plotSil(good_of_cluster$sil_k[[14]])
plotSil(good_of_cluster$sil_k[[15]])
plotSil(good_of_cluster$sil_k[[16]])
plotSil(good_of_cluster$sil_k[[17]])
plotSil(good_of_cluster$sil_k[[18]])
plotSil(good_of_cluster$sil_k[[19]])
k <- 16

# Perform k-means clustering with best k and compute the rand indices
set.seed(3)
km_out <- kmeans(absentData, k, nstart = 20)
km_clusters <- km_out$cluster
rand.index(km_clusters, as.numeric(as.factor(absentData_lab)))
adj.rand.index(km_clusters, as.numeric(as.factor(absentData_lab)))

# ----- HIERARCHICAL CLUSTERING AFTER PCA ----- #
# Proportion of variance explained
set.seed(3)
pr_out <- prcomp(absentData, scale = TRUE)
plot(pr_out)
```

```r
pve <- 100 * pr_out$sdev^2 / sum(pr_out$sdev^2)
plot(pve, type = "o",
xlab = "Principal Component", col = "blue", ylab = "Percent of variance explained (PVE)")

plot(cumsum(pve), type = "o", ylab = "Cumulative PVE",
     xlab = "Principal Component", col = "brown3")

# Get silhouette scores for multiple k values and plot them
set.seed(3)
k <- c(2:20)
good_of_cluster <- chooseK(k, hcSilK, pr_out$x[, 1:2])
good_of_cluster$plot

# Based off silhouette plots, choose the best k
plotSil(good_of_cluster$sil_k[[14]])
plotSil(good_of_cluster$sil_k[[15]])
plotSil(good_of_cluster$sil_k[[16]])
plotSil(good_of_cluster$sil_k[[17]])
plotSil(good_of_cluster$sil_k[[18]])
plotSil(good_of_cluster$sil_k[[19]])
k <- good_of_cluster$max_k

set.seed(3)
hc_out <- hclust(dist(dist(pr_out$x[, 1:2])))
hc_clusters <- cutree(hc_out, k)
rand.index(hc_clusters, as.numeric(as.factor(absentData_lab)))
adj.rand.index(hc_clusters, as.numeric(as.factor(absentData_lab)))
```

References

R Core Team. (2023). *R: A language and environment for statistical computing*. R Foundation for Statistical Computing. https://www.R-project.org/